



Deep Reinforcement Learning

Eligibility Traces

February 10, 2025

Vinay Kumar | Amberg

1. Problem Statement
2. Concept of Eligibility Traces
3. Modern Adaption
4. Results

How do we find good G_t ?

- Should the agent rely on immediate reward ?

How do we find good G_t ?

- Should the agent rely on immediate reward ?
- Or should agent consider long-term reward ?

TD Methods

- After taking action a_t , observe reward r_{t+1} and next state s_{t+1} .

TD Methods

- After taking action a_t , observe reward r_{t+1} and next state s_{t+1} .
- One-step TD target (G_t) and TD error (δ_t):

$$G_t = r_{t+1} + \gamma V(s_{t+1})$$

$$\delta_t = G_t - V(s_t) = r_{t+1} + \gamma V(s_{t+1}) - V(s_t)$$

where γ is the discount factor.

TD Methods

- After taking action a_t , observe reward r_{t+1} and next state s_{t+1} .
- One-step TD target (G_t) and TD error (δ_t):

$$G_t = r_{t+1} + \gamma V(s_{t+1})$$

$$\delta_t = G_t - V(s_t) = r_{t+1} + \gamma V(s_{t+1}) - V(s_t)$$

where γ is the discount factor.

- Update the value of the current state s_t :

$$V(s_t) \leftarrow V(s_t) + \alpha \delta_t$$

where α is the learning rate.

Monte Carlo Methods

- Play an episode until the end of terminal state (e.g., game over or goal reached).

Monte Carlo Methods

- Play an episode until the end of terminal state (e.g., game over or goal reached).
- Compute the total return G_t from time t to the end:

$$G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots$$

where γ is the discount factor.

Monte Carlo Methods

- Play an episode until the end of terminal state (e.g., game over or goal reached).
- Compute the total return G_t from time t to the end:

$$G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots$$

where γ is the discount factor.

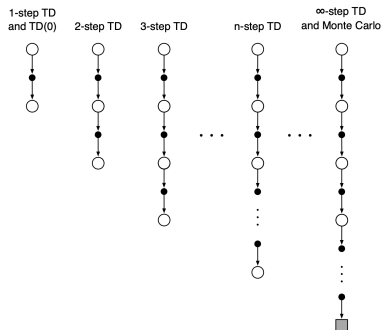
- For each state s_t visited in the episode, update its value:

$$V(s_t) \leftarrow V(s_t) + \alpha(G_t - V(s_t))$$

where α is the learning rate.

n-step TD

- **Idea:** Look Farther into future when do TD backup (1, 2, 3, ..., n steps)



<https://bjpcjp.github.io/pdfs/math/nstep-bootstrap-RL.pdf>

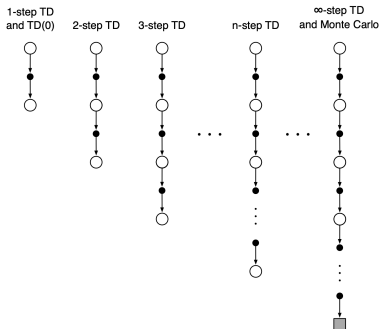
n-step TD

- **Idea:** Look Farther into future when do TD backup (1, 2, 3, ..., n steps)
- The n-step return $G_t^{(n)}$ is given by:

$$G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V(S_{t+n})$$



Where:

- $R_{t+1}, R_{t+2}, \dots, R_{t+n}$: rewards at each step
- γ : discount factor
- $V(S_{t+n})$: value of state at time $t + n$




<https://bjpcjp.github.io/pdfs/math/nstep-bootstrap-RL.pdf>

Action values increased by one-step Sarsa



How can we learn most from the episode?



<https://bjpcjp.github.io/pdfs/math/eligibility-traces-RL.pdf>

- A bridge from TD to Monte Carlo methods


- A bridge from TD to Monte Carlo methods
- A temporary record of which states/actions are eligible for learning updates

Key Concepts

The Forward View



- The idea is to create an average of our n-step




<http://incompleteideas.net/book/RLbook2020.pdf>- Chapter 12

Key Concepts

The Forward View

- The idea is to create an average of our n-step
- Weighting of past experiences in eligibility traces is controlled by λ (lambda)




<http://incompleteideas.net/book/RLbook2020.pdf>- Chapter 12

Key Concepts

The Forward View

- The idea is to create an average of our n-step
- Weighting of past experiences in eligibility traces is controlled by λ (lambda)
- Weights on the component returns are positive and sum to 1




<http://incompleteideas.net/book/RLbook2020.pdf>- Chapter 12

Key Concepts

The Forward View

- The idea is to create an average of our n-step
- Weighting of past experiences in eligibility traces is controlled by λ (lambda)
- Weights on the component returns are positive and sum to 1
- This particular way of averaging n-step called as TD(λ)




<http://incompleteideas.net/book/RLbook2020.pdf>- Chapter 12

Key Concepts

The Forward View

- The idea is to create an average of our n-step
- Weighting of past experiences in eligibility traces is controlled by λ (lambda)
- Weights on the component returns are positive and sum to 1
- This particular way of averaging n-step called as TD(λ)
- Each state has a trace that keeps track how much credit that state should receive for future updates.




<http://incompleteideas.net/book/RLbook2020.pdf>- Chapter 12

How weighting n sequence works ?

TD(λ)

- The one-step return is the largest weight, given by $1-\lambda$




<http://incompleteideas.net/book/RLbook2020.pdf>- Chapter 7

How weighting n sequence works ?

TD(λ)

- The one-step return is the largest weight, given by $1-\lambda$
- The weights fade away by λ with each additional steps




<http://incompleteideas.net/book/RLbook2020.pdf>- Chapter 7

How weighting n sequence works ?

TD(λ)

- The one-step return is the largest weight, given by $1-\lambda$
- The weights fade away by λ with each additional steps
- If ($\lambda = 0$) ?




<http://incompleteideas.net/book/RLbook2020.pdf> - Chapter 7

How weighting n sequence works ?

TD(λ)

- The one-step return is the largest weight, given by $1-\lambda$
- The weights fade away by λ with each additional steps
- If ($\lambda = 0$) ?
- Overall update reduces to its first component, one-step TD




<http://incompleteideas.net/book/RLbook2020.pdf> - Chapter 7

How weighting n sequence works ?

TD(λ)

- The one-step return is the largest weight, given by $1-\lambda$
- The weights fade away by λ with each additional steps
- If ($\lambda = 0$) ?
- Overall update reduces to its first component, one-step TD
- If ($\lambda = 1$) ?




<http://incompleteideas.net/book/RLbook2020.pdf>- Chapter 7


How weighting n sequence works ?

TD(λ)

- The one-step return is the largest weight, given by $1-\lambda$
- The weights fade away by λ with each additional steps
- If ($\lambda = 0$) ?
- Overall update reduces to its first component, one-step TD
- If ($\lambda = 1$) ?
- Overall update reduces to last component, i.e Monte Carlo




<http://incompleteideas.net/book/RLbook2020.pdf> - Chapter 7



Weighting given in the λ -return to each of the n -step returns.

- λ characterizes how fast the exponential weighting falls off


<https://bjpcjp.github.io/pdfs/math/eligibility-traces-RL.pdf>



Weighting given in the λ -return to each of the n -step returns.

- λ characterizes how fast the exponential weighting falls off
- Thus how far into the future the λ look like determining the weights

<https://bjpcjp.github.io/pdfs/math/eligibility-traces-RL.pdf>




Weighting given in the λ -return to each of the n-step returns.

- λ characterizes how fast the exponential weighting falls off
- Thus how far into the future the λ look like determining the weights
- Known as Off-line λ -return algorithm


<https://bjpcjp.github.io/pdfs/math/eligibility-traces-RL.pdf>

- Previously we discussed forward view (theoretical perspective)




<https://bjpcjp.github.io/pdfs/math/nstep-bootstrap-RL.pdf>

- Previously we discussed forward view (theoretical perspective)
- What we do is, we look ahead n steps for future rewards, and then we update our return




<https://bjpcjp.github.io/pdfs/math/nstep-bootstrap-RL.pdf>

- Previously we discussed forward view (theoretical perspective)
- What we do is, we look ahead n steps for future rewards, and then we update our return
- **Problem arises when:**




<https://bjpcjp.github.io/pdfs/math/nstep-bootstrap-RL.pdf>

- Previously we discussed forward view (theoretical perspective)
- What we do is, we look ahead n steps for future rewards, and then we update our return
- **Problem arises when:**
- Cannot update until future rewards are known




<https://bjpcjp.github.io/pdfs/math/nstep-bootstrap-RL.pdf>

- Previously we discussed forward view (theoretical perspective)
- What we do is, we look ahead n steps for future rewards, and then we update our return
- **Problem arises when:**
- Cannot update until future rewards are known
- Somehow complex to implement because the update of each state depends on later events or rewards that are not available at current time




<https://bjpcjp.github.io/pdfs/math/nstep-bootstrap-RL.pdf>

- Additional short term memory variable associated with each state




<https://bjpcjp.github.io/pdfs/math/nstep-bootstrap-RL.pdf>

- Additional short term memory variable associated with each state
- Updates the value function continuously




<https://bjpcjp.github.io/pdfs/math/nstep-bootstrap-RL.pdf>

- Additional short term memory variable associated with each state
- Updates the value function continuously
- When a reward is received, not just the current state, but also the past states get some credit



<https://bjpcjp.github.io/pdfs/math/nstep-bootstrap-RL.pdf>

- Additional short term memory variable associated with each state
- Updates the value function continuously
- When a reward is received, not just the current state, but also the past states get some credit
- Recently visited states get more credit, while older states get less credit over time.



<https://bjpcjp.github.io/pdfs/math/nstep-bootstrap-RL.pdf>

Semi-gradient TD(λ) for estimating $\hat{v} \approx v_\pi$

Input: the policy π to be evaluated

Input: a differentiable function $\hat{v} : \mathcal{S}^+ \times \mathbb{R}^d \rightarrow \mathbb{R}$ such that $\hat{v}(\text{terminal}, \cdot) = 0$

Algorithm parameters: step size $\alpha > 0$, trace decay rate $\lambda \in [0, 1]$

Initialize value-function weights \mathbf{w} arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$)

Loop for each episode:

 Initialize S

$\mathbf{z} \leftarrow \mathbf{0}$

(a d -dimensional vector)

 Loop for each step of episode:

 | Choose $A \sim \pi(\cdot|S)$

 | Take action A , observe R, S'

 | $\mathbf{z} \leftarrow \gamma \lambda \mathbf{z} + \nabla \hat{v}(S, \mathbf{w})$

 | $\delta \leftarrow R + \gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})$

 | $\mathbf{w} \leftarrow \mathbf{w} + \alpha \delta \mathbf{z}$


 | $S \leftarrow S'$

 until S' is terminal

<https://bjpcjp.github.io/pdfs/math/eligibility-traces-RL.pdf>

Comparison


Performance Evaluation



<https://bjpcjp.github.io/pdfs/math/eligibility-traces-RL.pdf>

Comparison

Performance Evaluation



<https://bjpcjp.github.io/pdfs/math/eligibility-traces-RL.pdf>

Generalized Advantage Estimation (GAE)

- The advantage function measures how much better (or worse) taking a specific action is compared to the average following the current policy in a given state

Generalized Advantage Estimation (GAE)

- The advantage function measures how much better (or worse) taking a specific action is compared to the average following the current policy in a given state
- The temporal difference (TD) error at each timestep is:

$$\delta_t^V = r_t + \gamma V(s_{t+1}) - V(s_t)$$

Generalized Advantage Estimation (GAE)

- The advantage function measures how much better (or worse) taking a specific action is compared to the average following the current policy in a given state
- The temporal difference (TD) error at each timestep is:

$$\delta_t^V = r_t + \gamma V(s_{t+1}) - V(s_t)$$

- The total return G_t used in advantage calculation:

$$G_t = \delta_t + \gamma \delta_{t+1} + \dots + \gamma^{n-1} \delta_{t+n-1}$$

which represents discounted sum of future rewards.

Improvement

The idea of eligibility traces

Generalized Advantage Estimation (GAE)

- Helps to estimate balance between immediate vs future rewards

Improvement

The idea of eligibility traces

Generalized Advantage Estimation (GAE)

- Helps to estimate balance between immediate vs future rewards
- Better decisions by looking at both short and long-term outcomes

The idea of eligibility traces

Generalized Advantage Estimation (GAE)

- Helps to estimate balance between immediate vs future rewards
- Better decisions by looking at both short and long-term outcomes
- GAE sums over multiple TD errors from multiple time steps using eligibility traces

The idea of eligibility traces


Generalized Advantage Estimation (GAE)

- Helps to estimate balance between immediate vs future rewards
- Better decisions by looking at both short and long-term outcomes
- GAE sums over multiple TD errors from multiple time steps using eligibility traces
- Estimate the advantage function, which measures how much better an action is compared to the average action in a given state

The idea of eligibility traces

Generalized Advantage Estimation (GAE)

- Helps to estimate balance between immediate vs future rewards
- Better decisions by looking at both short and long-term outcomes
- GAE sums over multiple TD errors from multiple time steps using eligibility traces
- Estimate the advantage function, which measures how much better an action is compared to the average action in a given state
- Works particularly well with deep neural networks and replay buffers.



- GAE act as a modern solution for complex environments

- GAE act as a modern solution for complex environments
- GAE inherits eligibility traces, that provide better stability

- GAE act as a modern solution for complex environments
- GAE inherits eligibility traces, that provide better stability
- Improved performance in delayed reward scenarios

- GAE act as a modern solution for complex environments
- GAE inherits eligibility traces, that provide better stability
- Improved performance in delayed reward scenarios
- Usage in real world scenarios like Robotics control, Game AI which helps deciding what to do in situations where each choice affects the next

- GAE act as a modern solution for complex environments
- GAE inherits eligibility traces, that provide better stability
- Improved performance in delayed reward scenarios
- Usage in real world scenarios like Robotics control, Game AI which helps deciding what to do in situations where each choice affects the next
- Combining GAE with advance models like PPO, make stable learning in many environments particularly real-world robotics applications.

End !

Thanks for paying attention !



https://makeameme.org/meme/haben-sie-noch#google_vignette